

# Conv2D Parameters and Default Values

Wednesday, October 10, 2018 10:54 AM

## Conv2D layer API Default Parameter Values

```
keras.layers.Conv2D(filters, kernel_size, strides=(1, 1), padding='valid', data_format=None, dilation_rate=(1, 1),
activation=None, use_bias=True, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=None,
bias_regularizer=None, activity_regularizer=None, kernel_constraint=None, bias_constraint=None)
```

## Explanation from Official Documentation

- filters**: Integer, the dimensionality of the output space (i.e. the number of output filters in the convolution).
- kernel\_size**: An integer or tuple/list of 2 integers, specifying the height and width of the 2D convolution window. Can be a single integer to specify the same value for all spatial dimensions.
- strides**: An integer or tuple/list of 2 integers, **specifying the strides of the convolution along the height and width**. Can be a single integer to specify the same value for all spatial dimensions. Specifying any stride value != 1 is incompatible with specifying any **dilation\_rate** value != 1.
- padding**: one of "valid" or "same" (case-insensitive). Note that "same" is slightly inconsistent across backends with **strides** != 1, as described [here](#)
- data\_format**: A string, one of "channels\_last" or "channels\_first". The ordering of the dimensions in the inputs. "channels\_last" corresponds to inputs with shape (batch, height, width, channels) while "channels\_first" corresponds to inputs with shape (batch, channels, height, width). It defaults to the **image\_data\_format** value found in your Keras config file at ~/.keras/keras.json. If you never set it, then it will be "channels\_last".
- dilation\_rate**: an integer or tuple/list of 2 integers, specifying the dilation rate to use for dilated convolution. Can be a single integer to specify the same value for all spatial dimensions. Currently, specifying any **dilation\_rate** value != 1 is incompatible with specifying any stride value != 1.
- activation**: Activation function to use (see [activations](#)). If you don't specify anything, no activation is applied (ie. "linear" activation:  $a(x) = x$ ).
- use\_bias**: Boolean, whether the layer uses a bias vector.
- kernel\_initializer**: Initializer for the **kernel** weights matrix (see [initializers](#)).
- bias\_initializer**: Initializer for the bias vector (see [initializers](#)).
- kernel\_regularizer**: Regularizer function applied to the **kernel** weights matrix (see [regularizer](#)).
- bias\_regularizer**: Regularizer function applied to the bias vector (see [regularizer](#)).
- activity\_regularizer**: Regularizer function applied to the output of the layer (its "activation"). (see [regularizer](#)).
- kernel\_constraint**: Constraint function applied to the kernel matrix (see [constraints](#)).
- bias\_constraint**: Constraint function applied to the bias vector (see [constraints](#)).

From <https://keras.io/layers/convolutional/>

## Filters

- "+": more weights, higher possible overfitting
- "-": less weights, underfitting, slowly converge

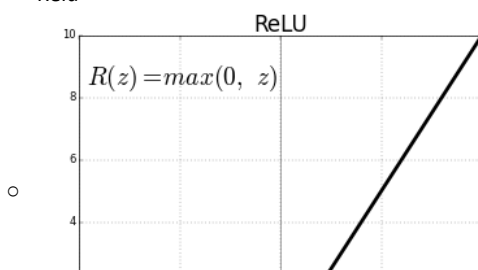
## Kernel\_size

Normally use 3, 5, 7

It's better to use small kernel size with more layers rather than large kernel size with less layers. The reason is related to **receptive field**.

## Activation functions

- Inside and outside the conv layers
- Choices:
  - Relu



## Kernel\_initializer: glorot\_normal

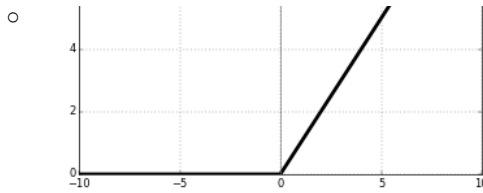
## Bias\_initializer: zeros

Choices:

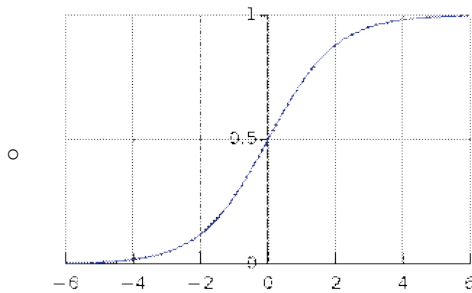
zeros, ones, constant, RandomNormal, RandomUniform, TruncatedNormal, VarianceScaling, Orthogonal, Identity, lecun\_uniform, glorot\_normal, glorot\_uniform, he\_normal, lecun\_normal, he\_uniform

## kernel\_regularizer, bias\_regularizer, activity\_regularizer

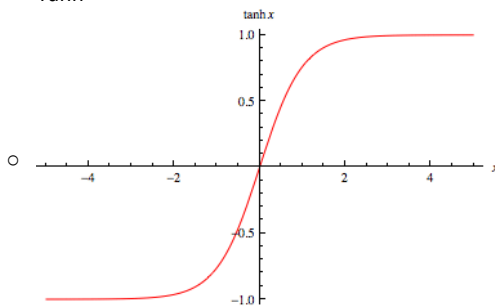
Regularizers allow to apply penalties on layer parameters or layer activity during



○ Sigmoid



○ Tanh



- Relu is a better choice than sigmoid or tanh  
This is because relu has two benefits:  
sparsity and lateral inhibition which can reduce the  
likelihood of vanishing gradient.

`kernel_regularizer, bias_regularizer, activity_regularizer`

Regularizers allow to apply penalties on layer parameters or layer activity during optimization.

When model get overfitted, use regularizer to improve the ability of generlization of your model.

Example:

```
from keras import regularizers
model.add(Dense(64, input_dim=64,
                  kernel_regularizer=regularizers.l2(0.01),
                  activity_regularizer=regularizers.l1(0.01)))
```

Available penalties:

```
keras.regularizers.l1(0.)
keras.regularizers.l2(0.)
keras.regularizers.l1_l2(l1=0.01, l2=0.01)
```

From <<https://keras.io/regularizers/>>

# Weights Initialization

Wednesday, October 10, 2018

10:34 AM

The normal vs uniform init seem to be rather unclear in fact.

If we refer solely on the [Glorot's](#) and [He's](#) initializations papers, they both use a similar theoretical analysis: they find a good variance for the distribution from which the initial parameters are drawn. This variance is adapted to the activation function used and is derived without explicitly considering the type of the distribution. As such, their theoretical conclusions hold for any type of distribution of the determined variance. In fact, in the Glorot paper, a uniform distribution is used whereas in the He paper it is a gaussian one that is chosen. The only "explanation" given for this choice in the He paper is:

Recent deep CNNs are mostly initialized by random weights drawn from Gaussian distributions

with a reference to [AlexNet paper](#). It was indeed released a little later than Glorot's initialization but however there is no justification in it of the use of a normal distribution.

In fact, in a [discussion on Keras issues tracker](#), they also seem to be a little confused and basically it could only be a matter of preference... (i.e. hypothetically Bengio would prefer uniform distribution whereas Hinton would prefer normal ones...) One the discussion, there is a small benchmark comparing Glorot initialization using a uniform and a gaussian distribution. In the end, it seems that the uniform wins but it is not really clear.

In the original [ResNet paper](#), it only says they used a gaussian He init for all the layers, I was not able to find where it is written that they used a uniform He init for the first layer. (maybe you could share a reference to this?)

As for the use of gaussian init with Batch Normalization, well, with BN the optimization process is less sensitive to initialization thus it is just a convention I would say.

From <<https://datascience.stackexchange.com/questions/13061/when-to-use-he-or-glorot-normal-initialization-over-uniform-init-and-what-are>>

# Ex1: Default model

Tuesday, October 9, 2018 11:02 AM

Ex1:

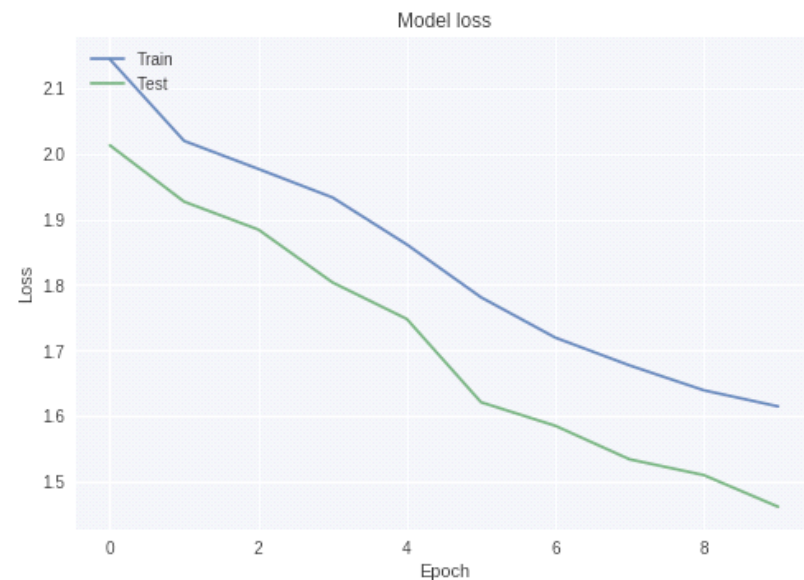
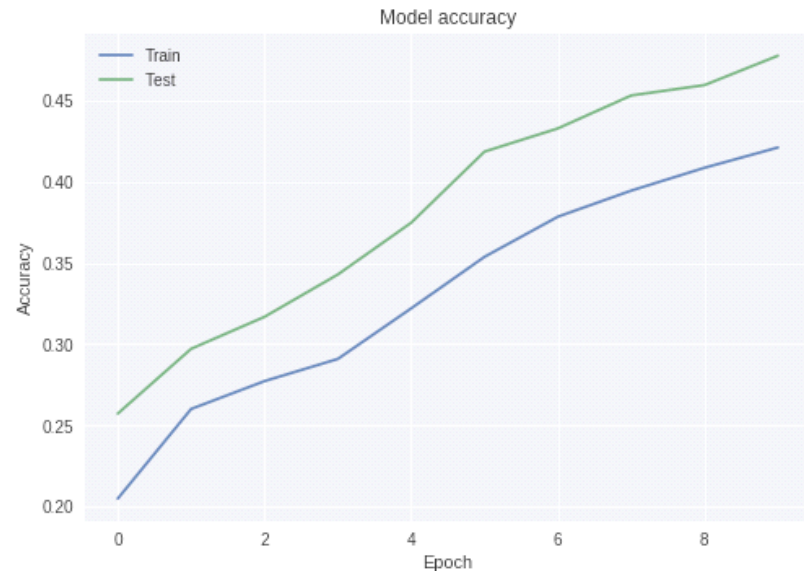
```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                input_shape=x_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

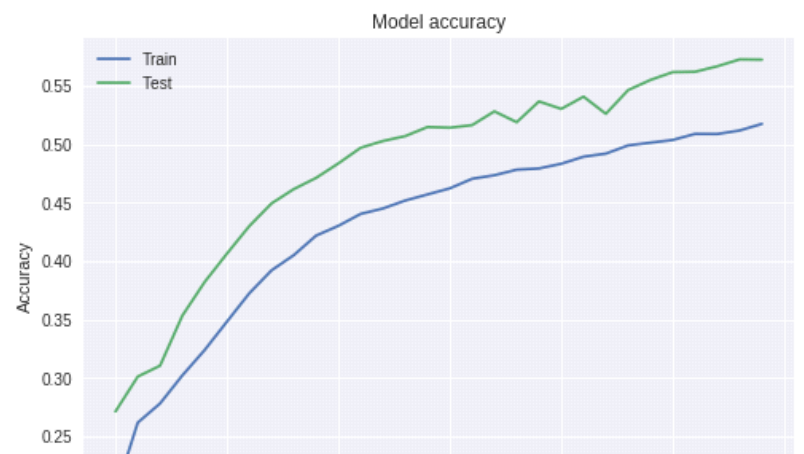
model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))
```

Epoch 1/100  
1563/1563 [=====] -  
52s 34ms/step - loss: 1.9137 - acc: 0.2947 - val\_loss:  
1.5951 - val\_acc: 0.4309  
Epoch 2/100  
1563/1563 [=====] -  
51s 33ms/step - loss: 1.6076 - acc: 0.4114 - val\_loss:  
1.5008 - val\_acc: 0.4652  
Epoch 3/100  
1563/1563 [=====] -  
51s 32ms/step - loss: 1.4872 - acc: 0.4602 - val\_loss:  
1.4133 - val\_acc: 0.5002  
Epoch 4/100  
1563/1563 [=====] -  
50s 32ms/step - loss: 1.4070 - acc: 0.4922 - val\_loss:  
1.2305 - val\_acc: 0.5558  
Epoch 5/100  
1563/1563 [=====] -  
50s 32ms/step - loss: 1.3472 - acc: 0.5173 - val\_loss:  
1.2247 - val\_acc: 0.5685

10 epochs



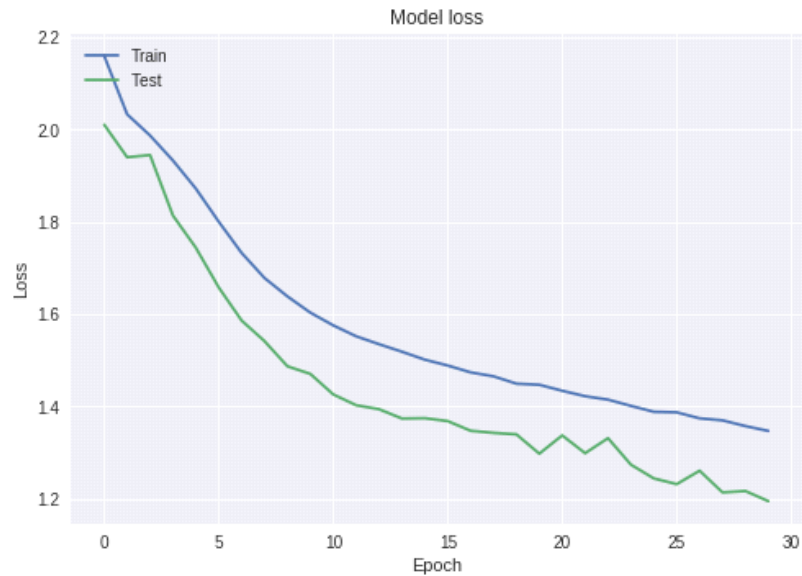
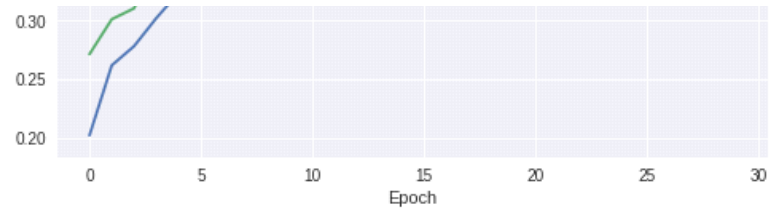
30 epochs



```

Epoch 5/100
1563/1563 [=====] -
50s 32ms/step - loss: 1.3472 - acc: 0.5173 - val_loss:
1.2242 - val_acc: 0.5685
Epoch 6/100
1563/1563 [=====] -
51s 32ms/step - loss: 1.2886 - acc: 0.5405 - val_loss:
1.1495 - val_acc: 0.5881
Epoch 7/100
1563/1563 [=====] -
51s 32ms/step - loss: 1.2314 - acc: 0.5633 - val_loss:
1.0979 - val_acc: 0.6164
Epoch 8/100
1563/1563 [=====] -
51s 33ms/step - loss: 1.1890 - acc: 0.5795 - val_loss:
1.0800 - val_acc: 0.6165
Epoch 9/100
1563/1563 [=====] -
51s 33ms/step - loss: 1.1551 - acc: 0.5907 - val_loss:
1.0318 - val_acc: 0.6294
Epoch 10/100
1563/1563 [=====] -
51s 32ms/step - loss: 1.1223 - acc: 0.6037 - val_loss:
0.9952 - val_acc: 0.6482
Epoch 11/100
1563/1563 [=====] -
51s 32ms/step - loss: 1.0910 - acc: 0.6168 - val_loss:
0.9556 - val_acc: 0.6660
Epoch 12/100
1563/1563 [=====] -
50s 32ms/step - loss: 1.0642 - acc: 0.6260 - val_loss:
0.9937 - val_acc: 0.6525
Epoch 13/100
1563/1563 [=====] -
51s 32ms/step - loss: 1.0439 - acc: 0.6309 - val_loss:
0.9177 - val_acc: 0.6807
Epoch 14/100
1563/1563 [=====] -
50s 32ms/step - loss: 1.0178 - acc: 0.6425 - val_loss:
0.9341 - val_acc: 0.6800
Epoch 15/100
1563/1563 [=====] -
51s 33ms/step - loss: 1.0001 - acc: 0.6506 - val_loss:
0.9134 - val_acc: 0.6855
Epoch 16/100
1563/1563 [=====] -
51s 33ms/step - loss: 0.9794 - acc: 0.6579 - val_loss:
0.8992 - val_acc: 0.6867
Epoch 17/100
1563/1563 [=====] -
50s 32ms/step - loss: 0.9590 - acc: 0.6652 - val_loss:
0.8850 - val_acc: 0.6964
Epoch 18/100
1563/1563 [=====] -
51s 32ms/step - loss: 0.9494 - acc: 0.6677 - val_loss:
0.8405 - val_acc: 0.7124
Epoch 19/100
1563/1563 [=====] -
51s 33ms/step - loss: 0.9361 - acc: 0.6763 - val_loss:
0.8200 - val_acc: 0.7175
Epoch 20/100
1563/1563 [=====] -
51s 33ms/step - loss: 0.9249 - acc: 0.6787 - val_loss:

```



0.8467 - val\_acc: 0.7037  
Epoch 21/100  
1563/1563 [=====] -  
51s 33ms/step - loss: 0.9189 - acc: 0.6815 - val\_loss:  
0.7741 - val\_acc: 0.7367

## Ex2: Double the numbers of filters

Tuesday, October 9, 2018 11:27 AM

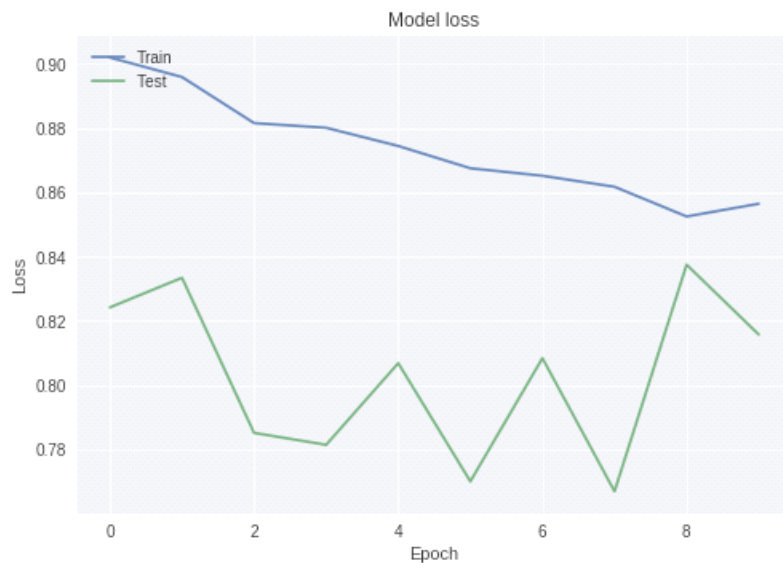
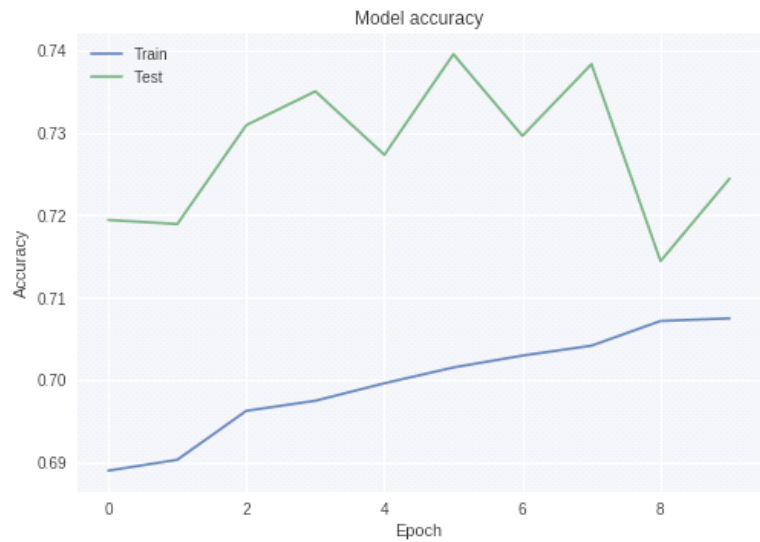
### Double the numbers of filters

Ex2:

```
model = Sequential()
model.add(Conv2D(64, (3, 3), padding='same',
                input_shape=x_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(128, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))
```



## Ex3: Inside Relu

Tuesday, October 9, 2018 11:51 AM

### Use the activation functions inside conv2D layers

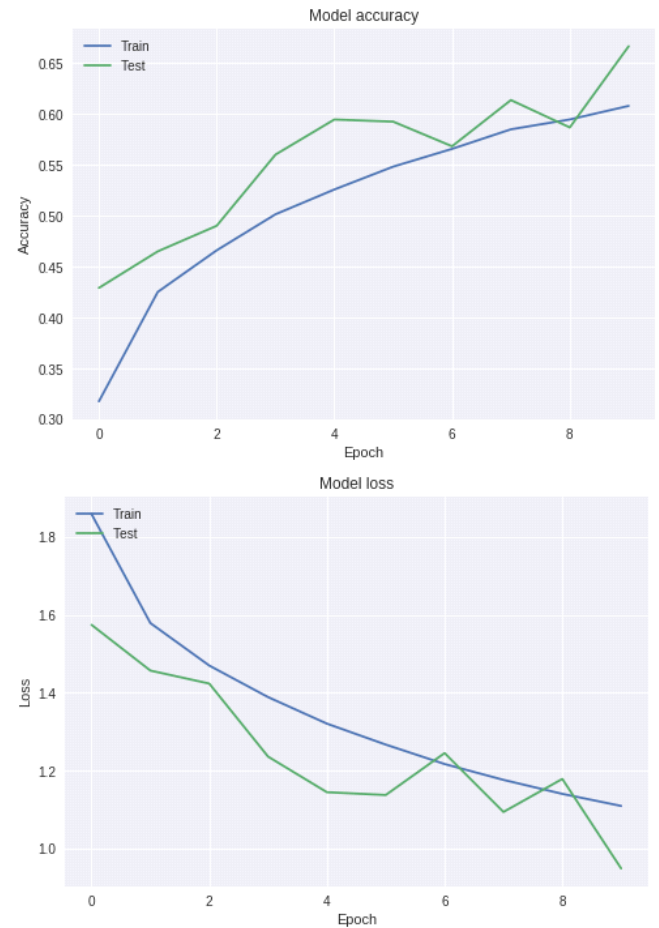
```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                input_shape=x_train.shape[1:], activation='relu'))
#model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3), activation='relu'))
#model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
#model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3), activation='relu'))
#model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

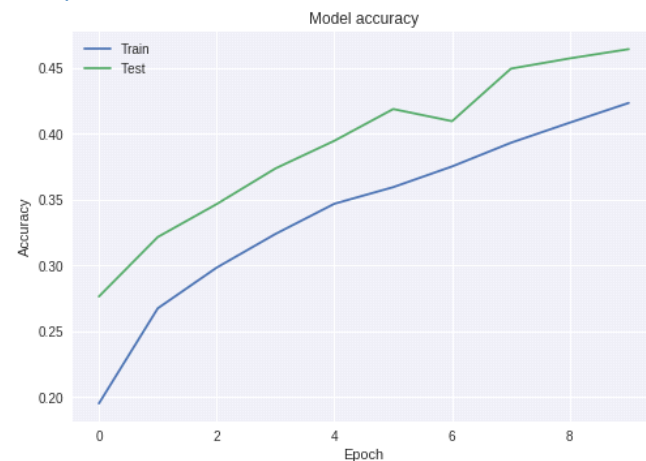
model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

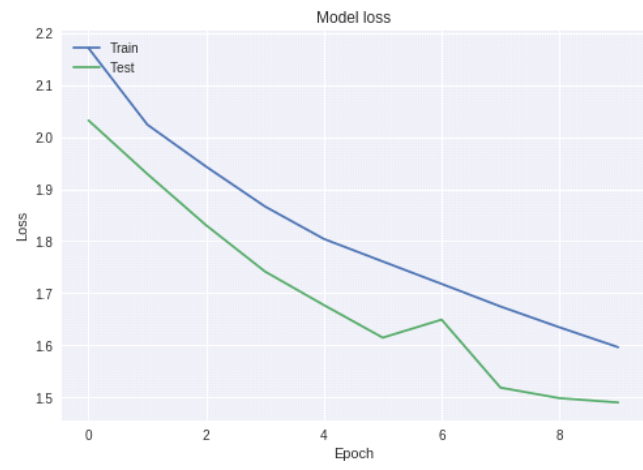
### 10 epochs



### 30 epochs







# Ex4: Not using bias inside conv2D layers

Tuesday, October 9, 2018 1:24 PM

## Not using bias inside conv2D layers

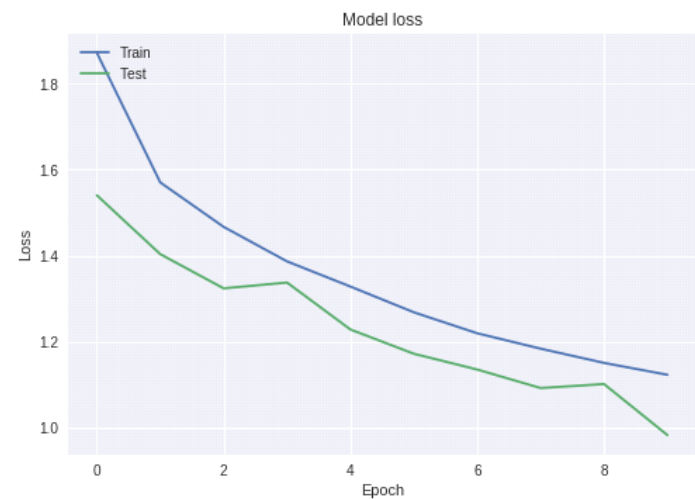
```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                 input_shape=x_train.shape[1:], use_bias=False))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3), use_bias=False))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same', use_bias=False))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3), use_bias=False))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

## 10 epochs



## Ex5: Using half numbers of filters

Tuesday, October 9, 2018 1:45 PM

### With less filters

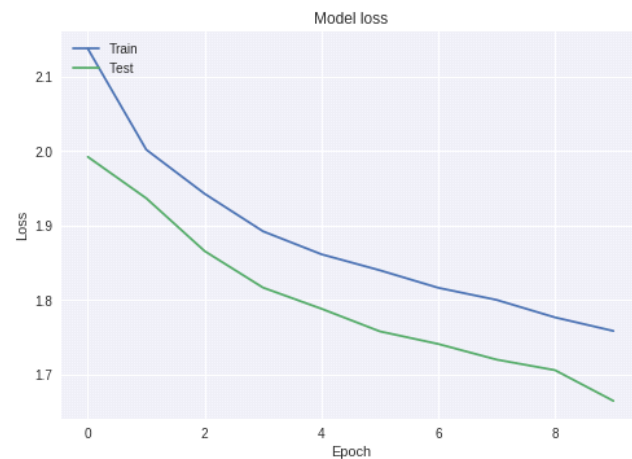
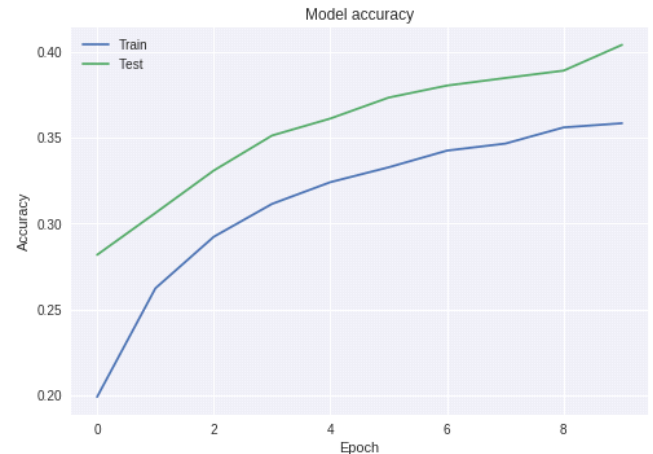
```
model = Sequential()
model.add(Conv2D(16, (3, 3), padding='same',
                input_shape=x_train.shape[1:], use_bias=False))
model.add(Activation('relu'))
model.add(Conv2D(16, (3, 3), use_bias=False))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(32, (3, 3), padding='same', use_bias=False))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3), use_bias=False))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

### 10 epochs



# Ex6: Triple the numbers of filters

Wednesday, October 10, 2018 10:21 AM

## Triple the numbers of filters

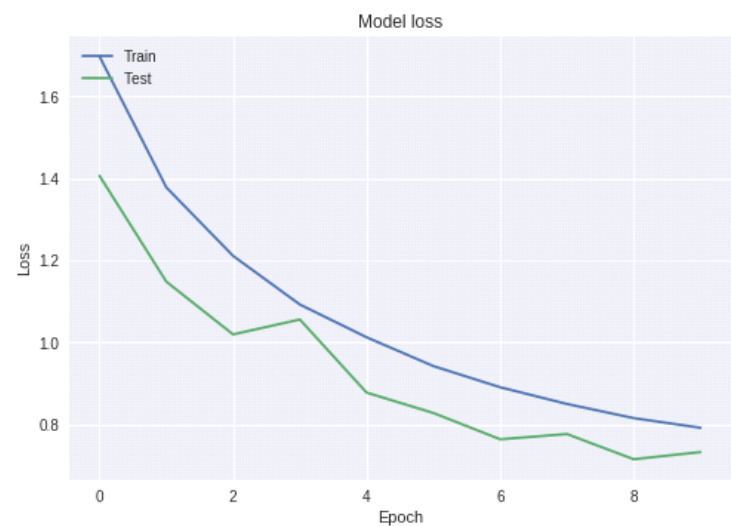
```
model = Sequential()
model.add(Conv2D(96, (3, 3), padding='same',
                input_shape=x_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(96, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(192, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(192, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

## 10 epochs



# Ex7: Using He normal

Wednesday, October 10, 2018 10:35 AM

## Using He normal

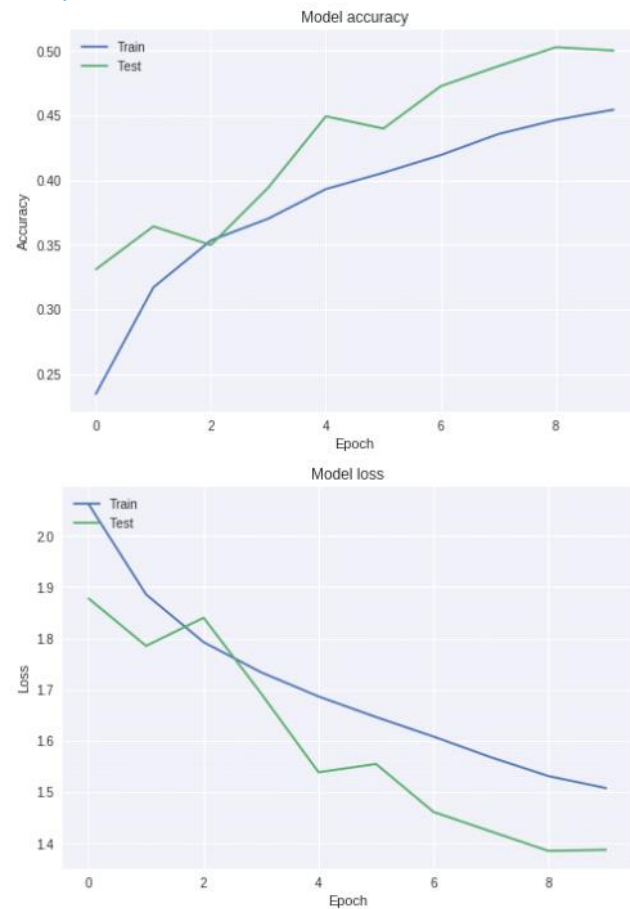
```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same', kernel_initializer='he_normal',
                input_shape=x_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3), kernel_initializer='he_normal'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same', kernel_initializer='he_normal'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3), kernel_initializer='he_normal'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

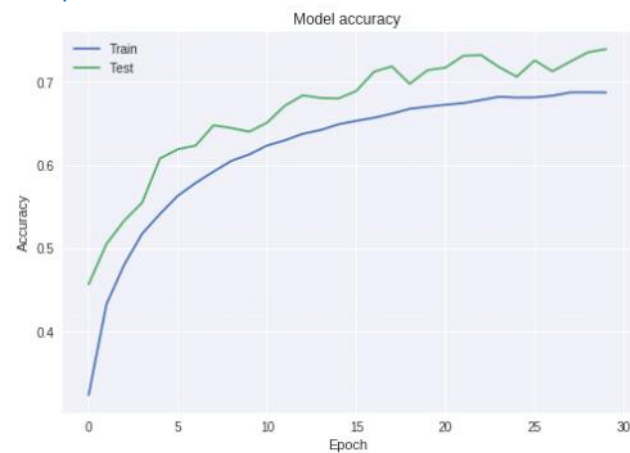
model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

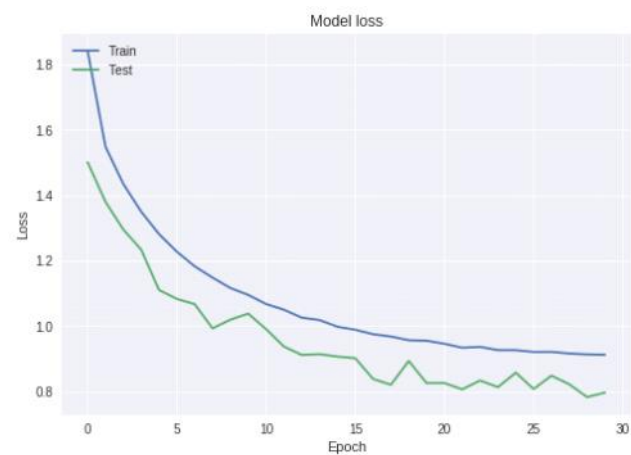
model.summary()
```

## 10 epochs



## 30 epochs





## Ex8: Inside Sigmoid

Wednesday, October 10, 2018 10:39 AM

```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                 input_shape=x_train.shape[1:], activation='sigmoid'))
model.add(Conv2D(32, (3, 3), activation='sigmoid'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same', activation='sigmoid'))
model.add(Conv2D(64, (3, 3), activation='sigmoid'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

30 epochs



## Ex9: Inside tanh

Wednesday, October 10, 2018 10:39 AM

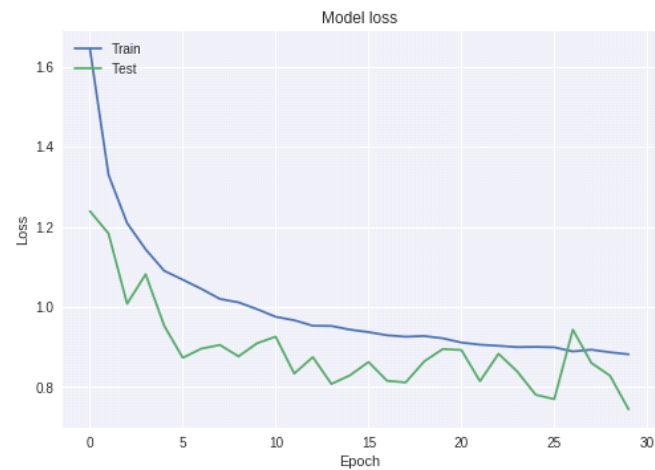
```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                 input_shape=x_train.shape[1:], activation='tan
model.add(Conv2D(32, (3, 3), activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same', activation='tan
model.add(Conv2D(64, (3, 3), activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

30 epochs





## Ex10: Using (5,5) kernel size

Wednesday, October 10, 2018 1:15 PM

```
model = Sequential()
model.add(Conv2D(32, (5, 5), padding='same',
                input_shape=x_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(32, (5, 5)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (5, 5), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(64, (5, 5)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

### 30 epochs

